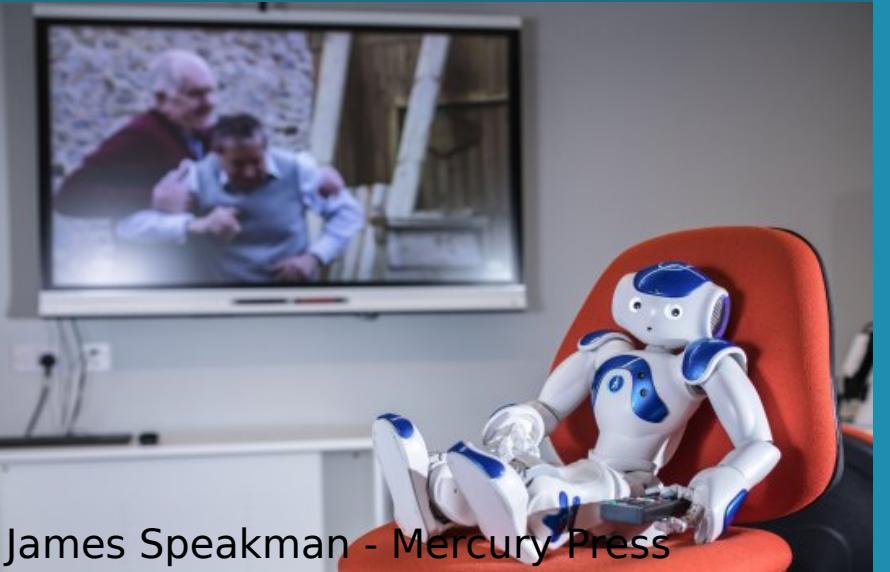


Introduction

Our goal is to learn **continually** from an infinite stream of data
without the constraint of **storing** the previous data
without catastrophic **forgetting** due to gradual data distribution shifts
without relying on **task boundaries**



Method

Memory Aware Synapses - method for task-incremental learning

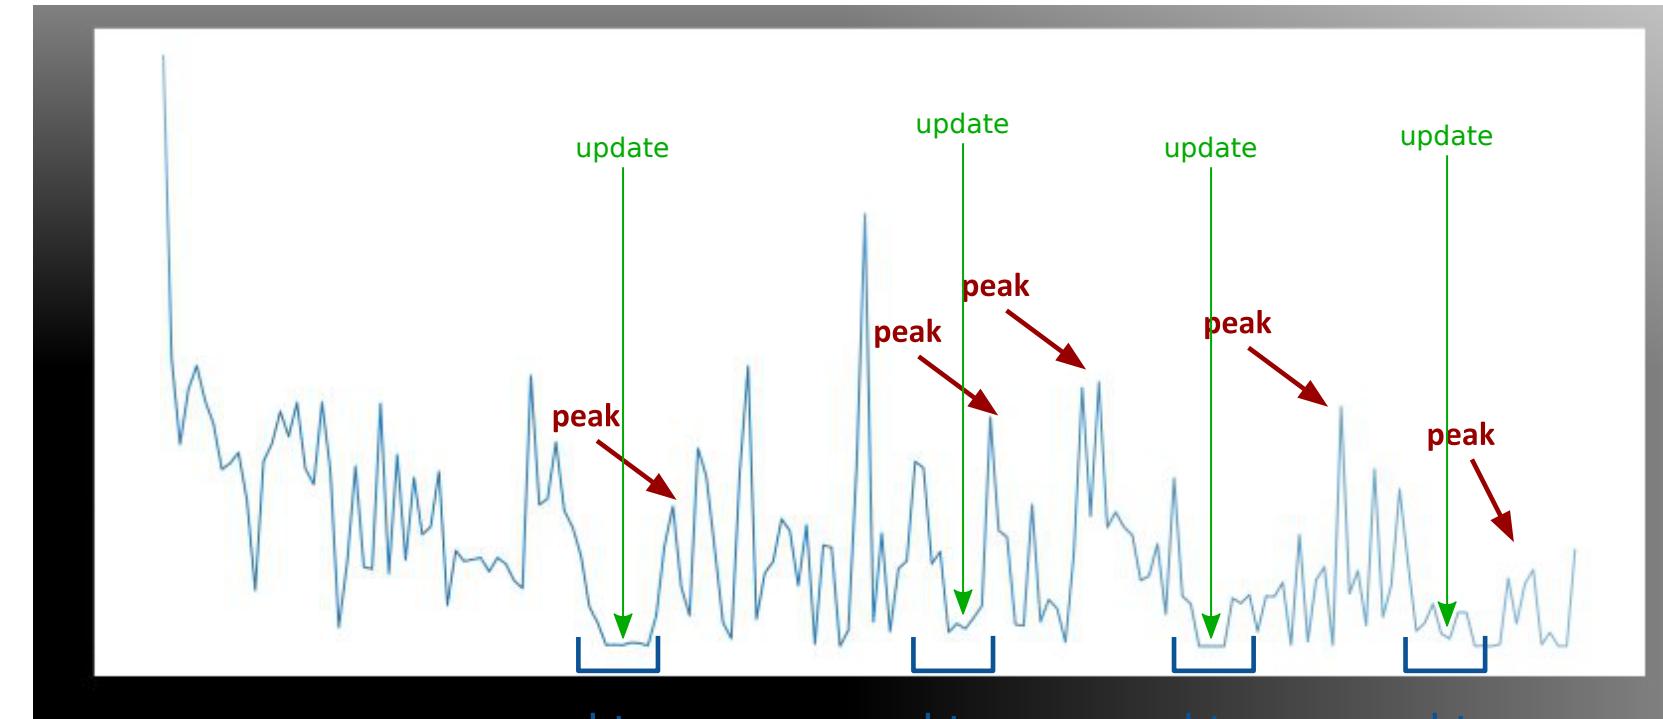
While training a new task, changes on important parameters are penalized:

$$F(x_k; \theta + \delta) - F(x_k; \theta) \approx \sum_i g_i(x_k) \delta_i \quad \Omega_i = \frac{1}{N} \sum_{k=1}^N \|g_i(x_k)\| \quad g_i(x_k) = \frac{\partial F(x_k)}{\partial \theta_i}$$

When training the new task, deviations are punished according to their importance:

$$L(\theta) = L_n(\theta) + \frac{\lambda}{2} \sum_i \Omega_i (\theta_i - \theta_i^*)^2$$

1) When to update importance weights?



Loss plateaus are stable moments to update the importance weights, once per plateau.

2) Which data to update them?

Small buffer with hard samples, B , (size 100) is used to stabilize the online setting. Hardest samples are selected.

3) How to accumulate them?

Instead of adding the importance weights, each update computes a **moving average**.

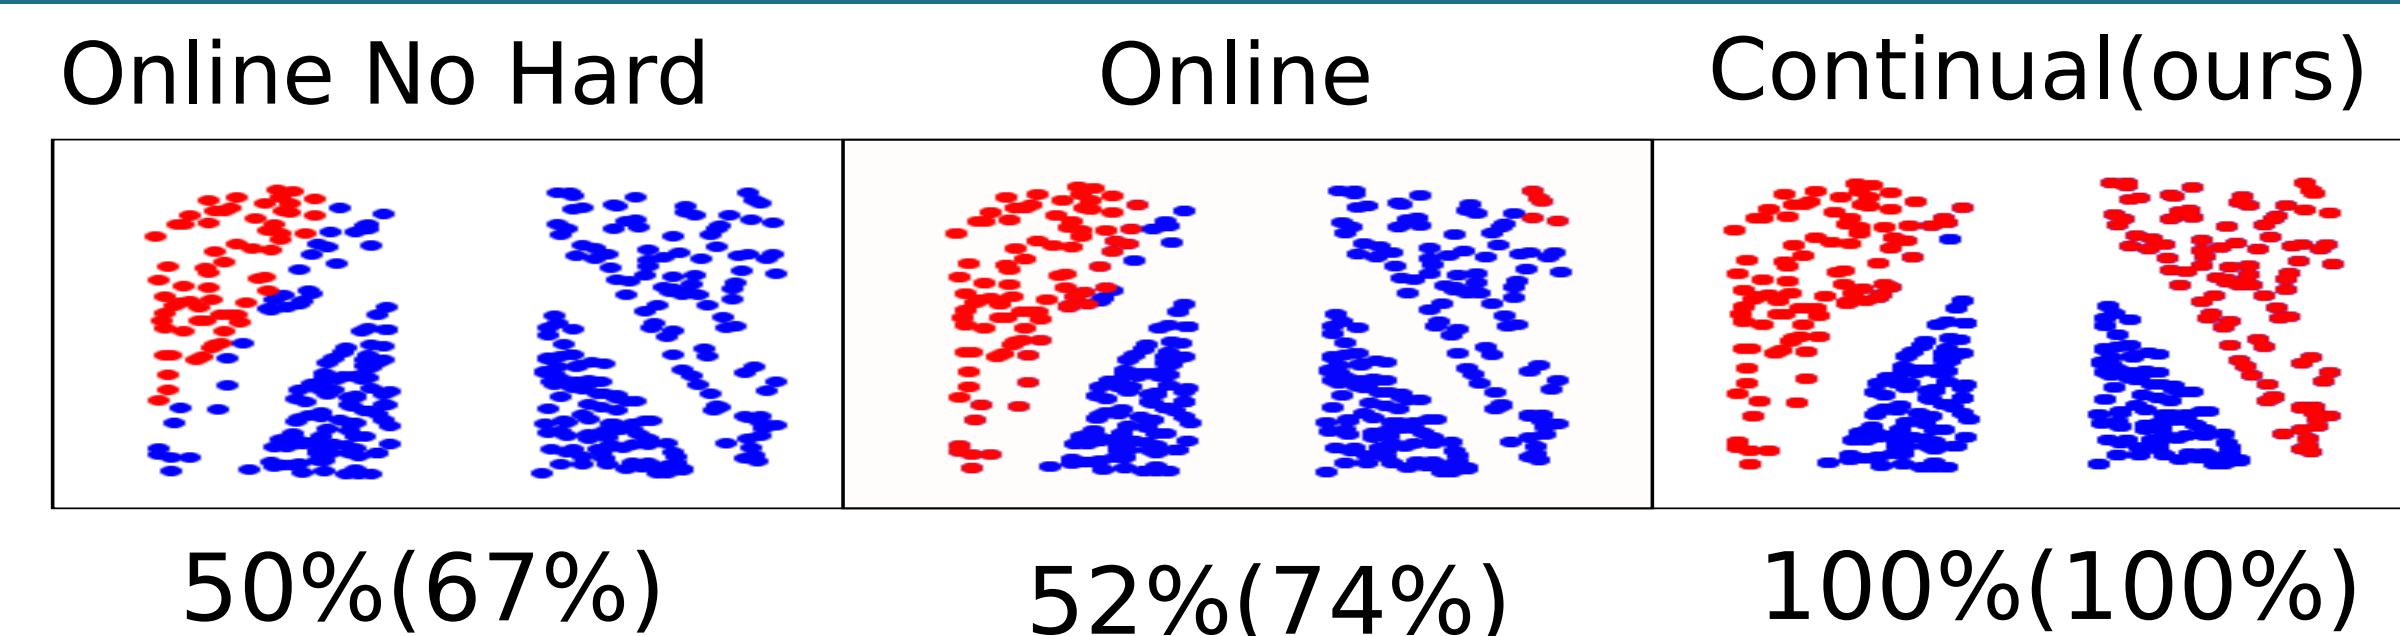
Algorithm 1 Online Continual Learning

```

1: Input:  $\delta_\mu, \delta_\sigma, \mathcal{N}$ 
2: Initialize:  $\mathcal{B} = \{\}, \mathcal{W} = \{\},$ 
3:  $\Omega = \vec{0}, \mu_L^{old} = 0, \sigma_L^{old} = 0, \mathcal{P} = 0$ 
4: repeat
5:   Receive  $K$  recent samples  $X, Y$ 
6:   for  $n$  in  $\mathcal{N}$  do
7:      $\mathcal{L}^T = \mathcal{L}_\theta(X, Y) + \mathcal{L}_\theta(X_B, Y_B) + \frac{\lambda}{2} \sum_i \Omega_i (\theta_i - \theta_i^*)^2$ 
8:      $\theta \leftarrow \text{SGD}(\theta, \mathcal{L}^T)$ 
9:     if  $n = 1$  then
10:        $\mathcal{W} \leftarrow \text{update}(\mathcal{W}, \mathcal{L}(X, Y), \mathcal{L}(X_B, Y_B))$ 
11:     end if
12:   end for
13:   if  $\neg \mathcal{P} \wedge \mu(\mathcal{W}) < \delta_\mu \wedge \sigma(\mathcal{W}) < \delta_\sigma$  then
14:      $\Omega \leftarrow \text{update}(\Omega, \theta, (X_B, Y_B))$ 
15:      $\mu_L^{old} = \mu(\mathcal{W}), \sigma_L^{old} = \sigma(\mathcal{W})$ 
16:      $\mathcal{W} = \{\}, \mathcal{P} = 1$ 
17:   end if
18:   if  $\mu(\mathcal{W}) > \mu_L^{old} + \sigma_L^{old}$  then
19:      $\mathcal{P} = 0$ 
20:   end if
21:    $(X_B, Y_B) \leftarrow \text{update}((X_B, Y_B), (X, Y), \mathcal{L}_\theta(X, Y))$ 

```

Synthetic experiment



- Classify points in or out of the unit circle.
- Data sampled from 2 quadrants representing task 1 and 2.
- Accuracy on first quadrant (and total performance).

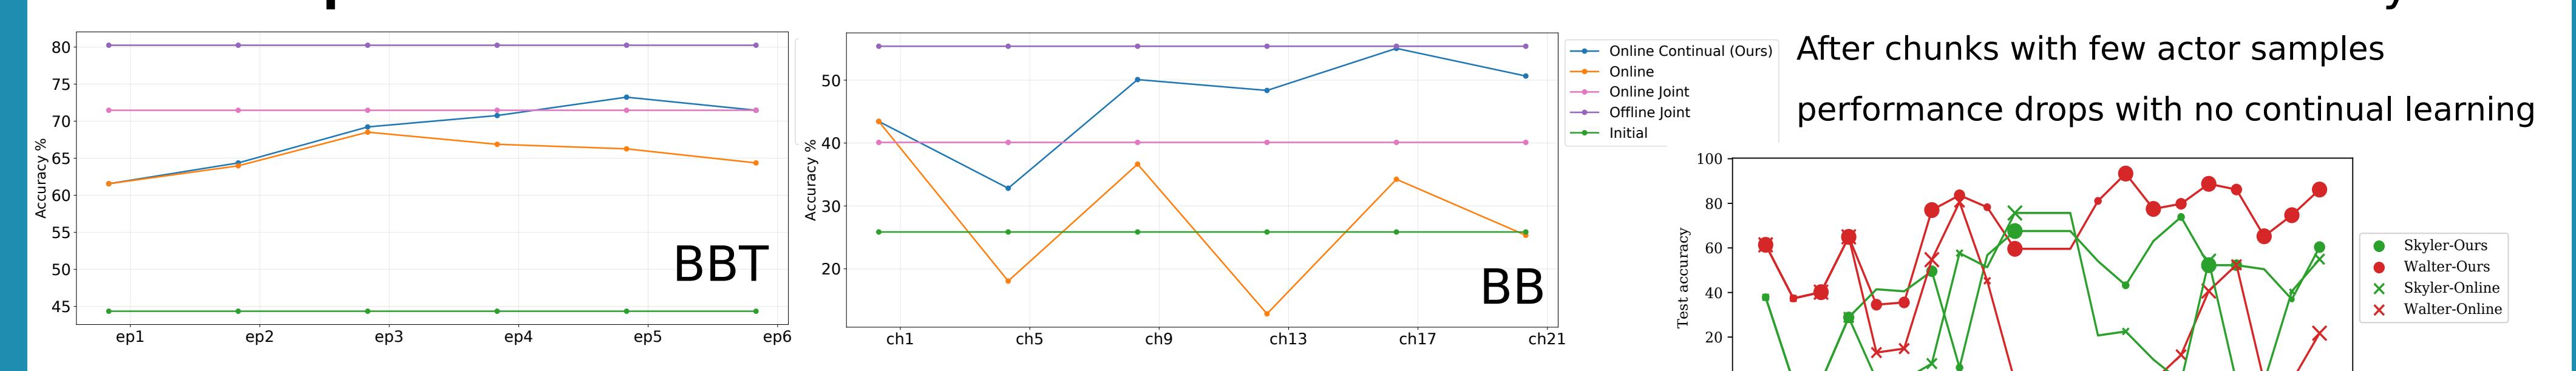
Face recognition in soap series

Goal: recognize faces of different actors.

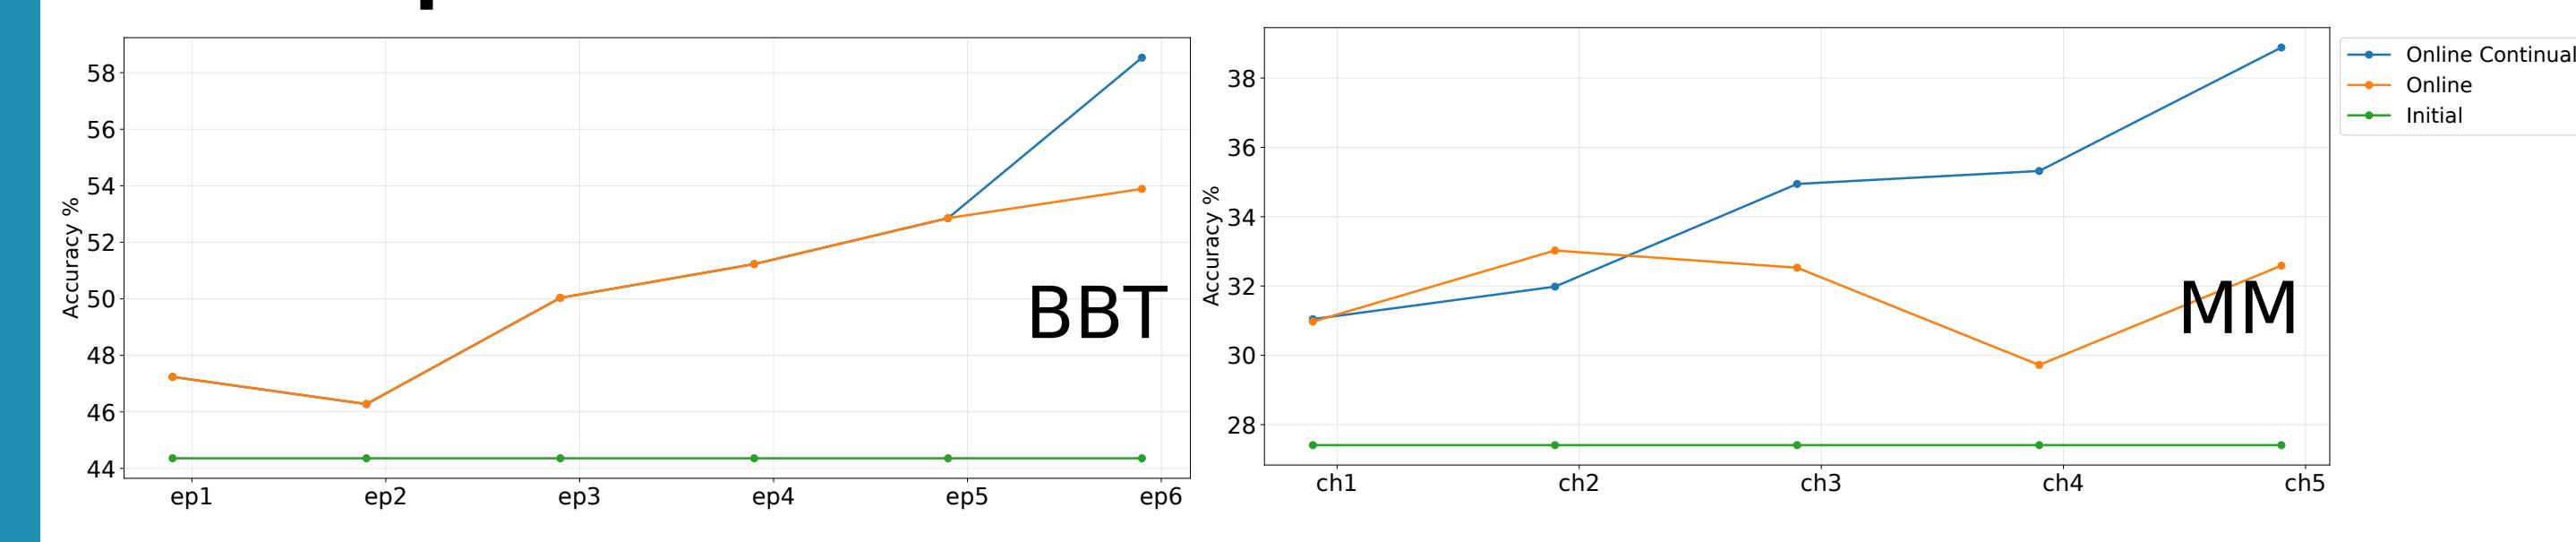
How:
using an online face detector module and a multi-object face tracker.
Actors are classified according to the distance to their closest template within an embedding space.



Weak Supervision: an annotator tells whether two consecutive tracks have the same identity.



Self Supervision: assume tracks of two faces from same frame must belong to different actors.



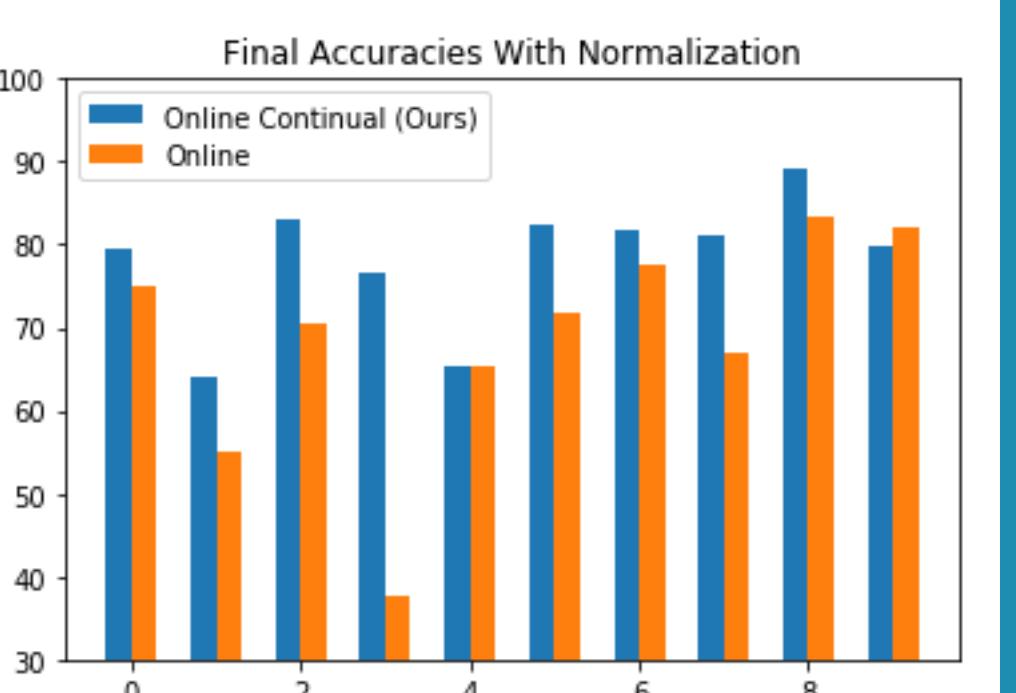
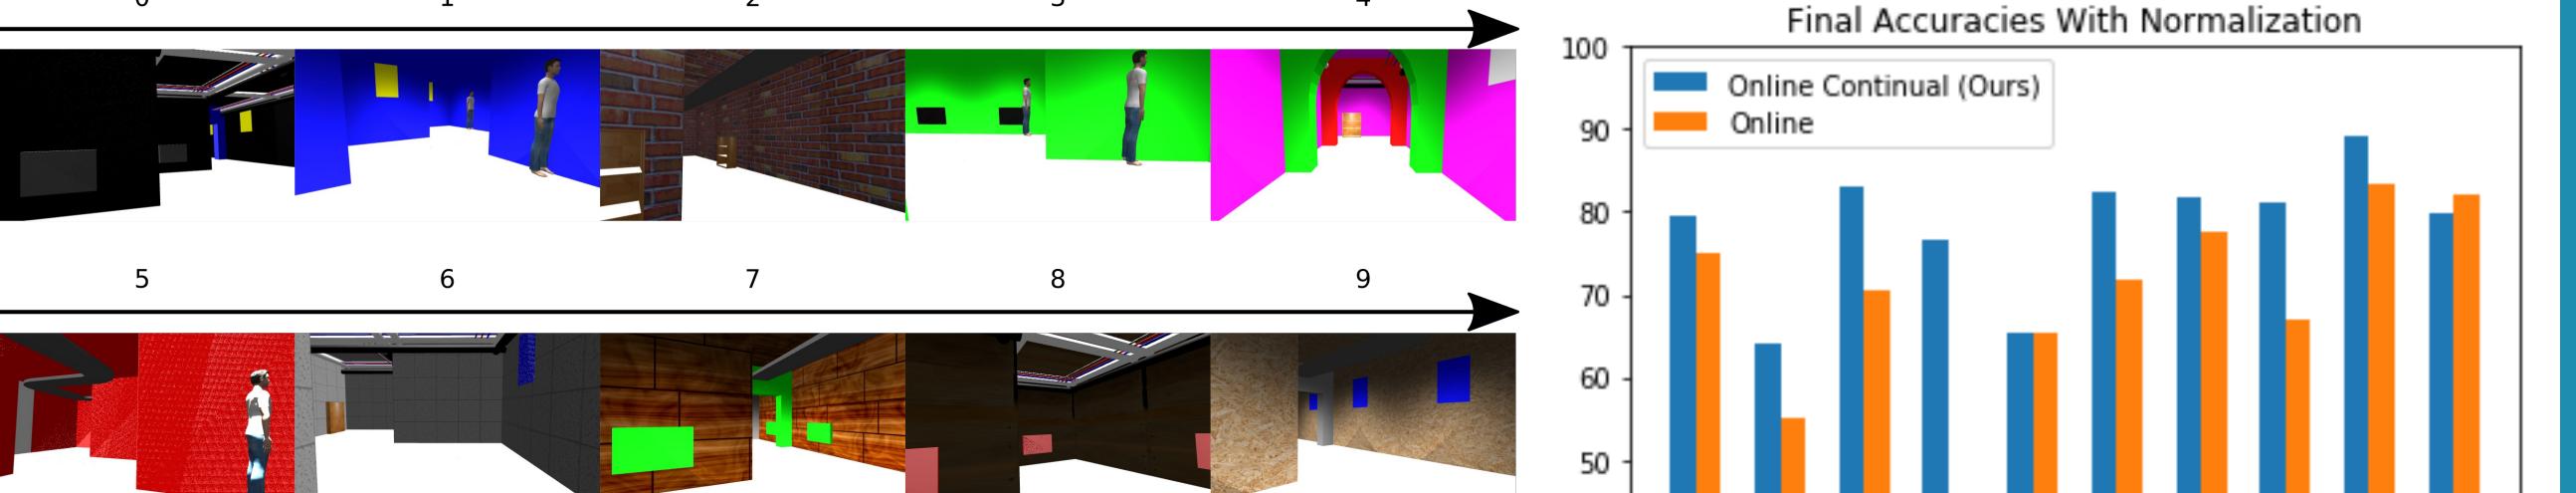
Monocular collision avoidance

Drone in ROS-Gazebo simulation

using off-policy simulation-supervision.

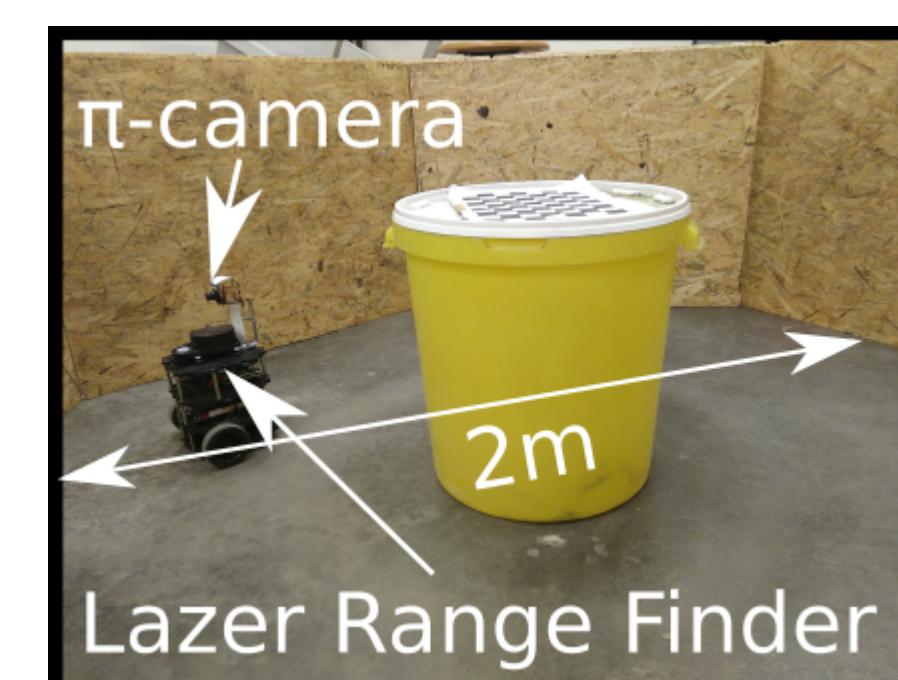
Depth based heuristic demonstrates navigation in sequence of 10 corridor.

Continual learning avoids forgetting flight in earlier corridors.



Turtlebot in the real world:

- Supervision from heuristic with range finder
- On-policy: agent controls while training (harder)
- Gradual changes in environment.



Output normalization:

sample experiences to have an equal distribution over actions.

Conclusion

Online importance weight-based **continual learning** by defining **when**, **on which data** and **how** to update importance weights.

Pushing boundaries from task-based to **task-free** and **online continual learning**. Method successfully validated on both **face recognition** and monocular **collision avoidance**.

Code: https://github.com/kkelchte/task_free_continual_learning